

# Location-dependent Query Processing

- Location-dependent applications
- Location-dependent queries
- Mobile positioning techniques
- Location-dependent databases
  - Data representation
  - Indexing and retrieval
  - Query processing

# Location-dependent applications

## Traffic Monitoring

- How many cars are in the downtown area?
- Send an alert if a non-friendly vehicle enters a restricted region
- Report any congestion in the road network
- Once an accident is discovered, immediately send alarm to the nearest police and ambulance cars
- Make sure that there are no two aircrafts with nearby paths

## Location-based Object Finder / Advertisement

- Where is my nearest Gas station?
- What are the fast food restaurants within 3 miles from my location?
- Am I near to a restaurant while any of my friends are there
- Send E-coupons to all customers within 3 miles of my stores
- Where is Bus 973 now? How long do I have to wait for its arrival?

# Location-dependent applications

## ■ Location-dependent applications have three key elements

- ◆ *Location-dependent queries:*

- ▶ Results depend on the location of the queries or the objects being queried for.

- ◆ *Locations:*

- ▶ Obtained by using mobile positioning technologies.
  - ▶ Maintained by spatial databases.

- ◆ *Location-dependent database:*

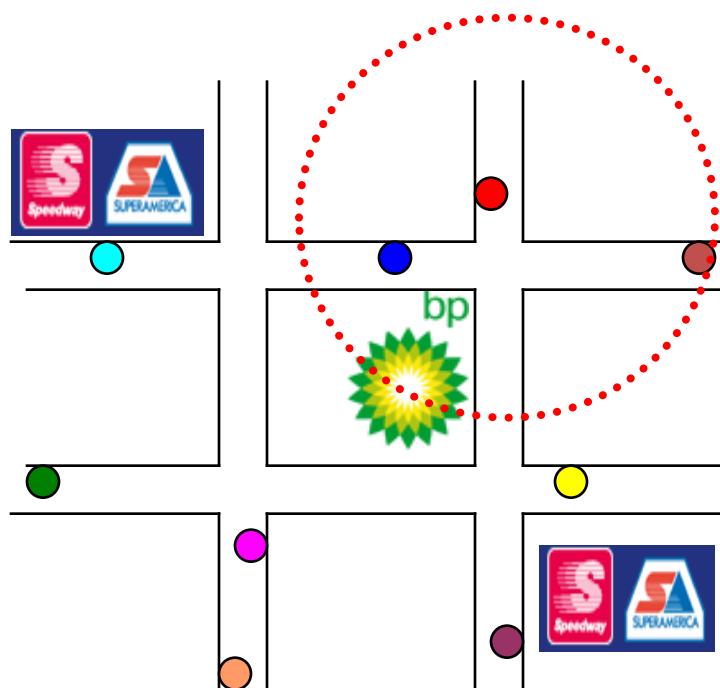
- ▶ Representation, storage, and retrieval of location-dependent data
  - ▶ Processing of location-dependent queries

# Location dependant query

- **Query whose result depends on the geographical location of the origin of the query**
  - ◆ Moving query
  - ◆ Example: What is the distance of Hung Hom station from here?
- **Query whose result depends on the geographical location of the object being queried for**
  - ◆ Moving object
  - ◆ Example: Where is the Taxi with plate no. 12345678?
- **User continuously asks the same question - every time the answer is different but correct.**
- **How about query whose context is related to locations?**
  - ◆ Find out the top 5 roads in KLW that have the most traffic between 7:00-8:00pm

# Location-dependent queries

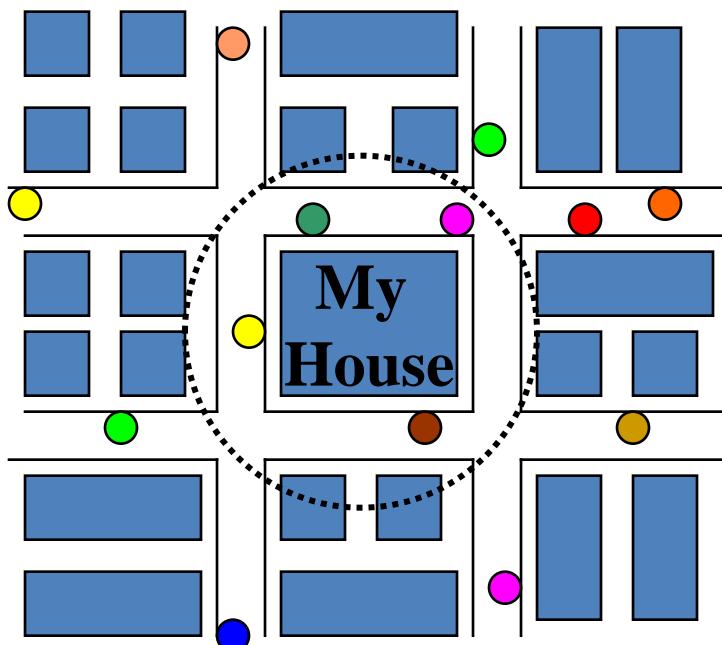
*Moving query on stationary objects*



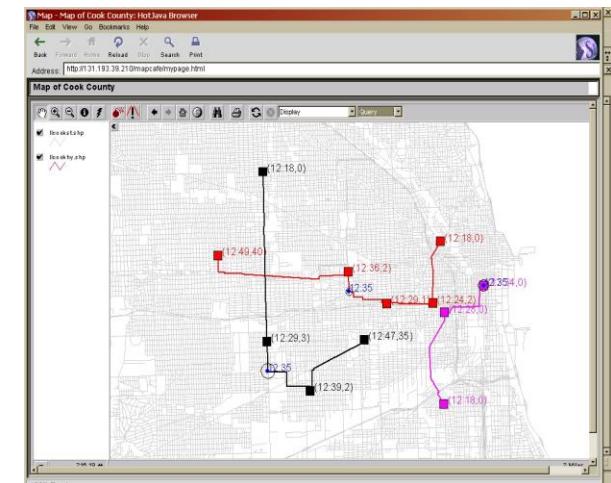
**Find the nearest gas station(s) within 1 miles of moving red car**

# Location-dependent queries

*Stationary query on moving objects*

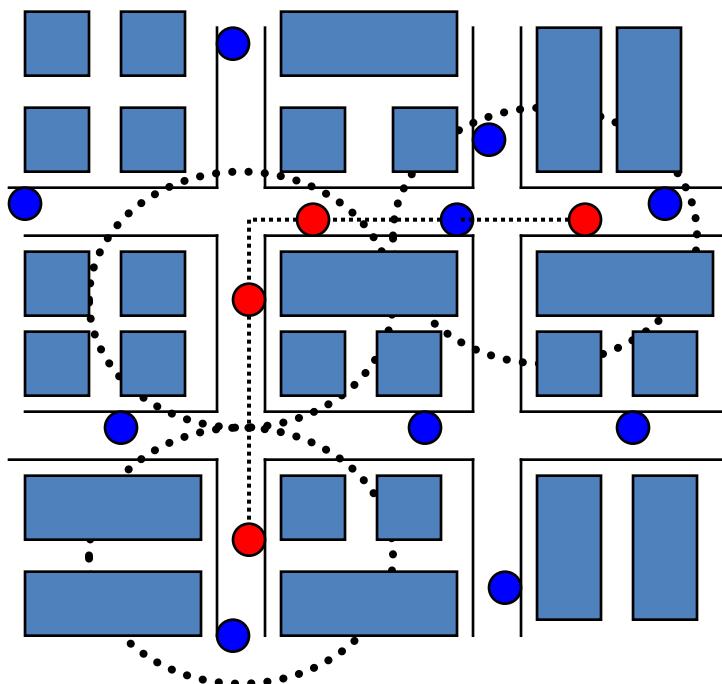


**Continuously find all vehicles within 1 miles of my house**



# Location-dependent queries

*Moving query on moving objects*

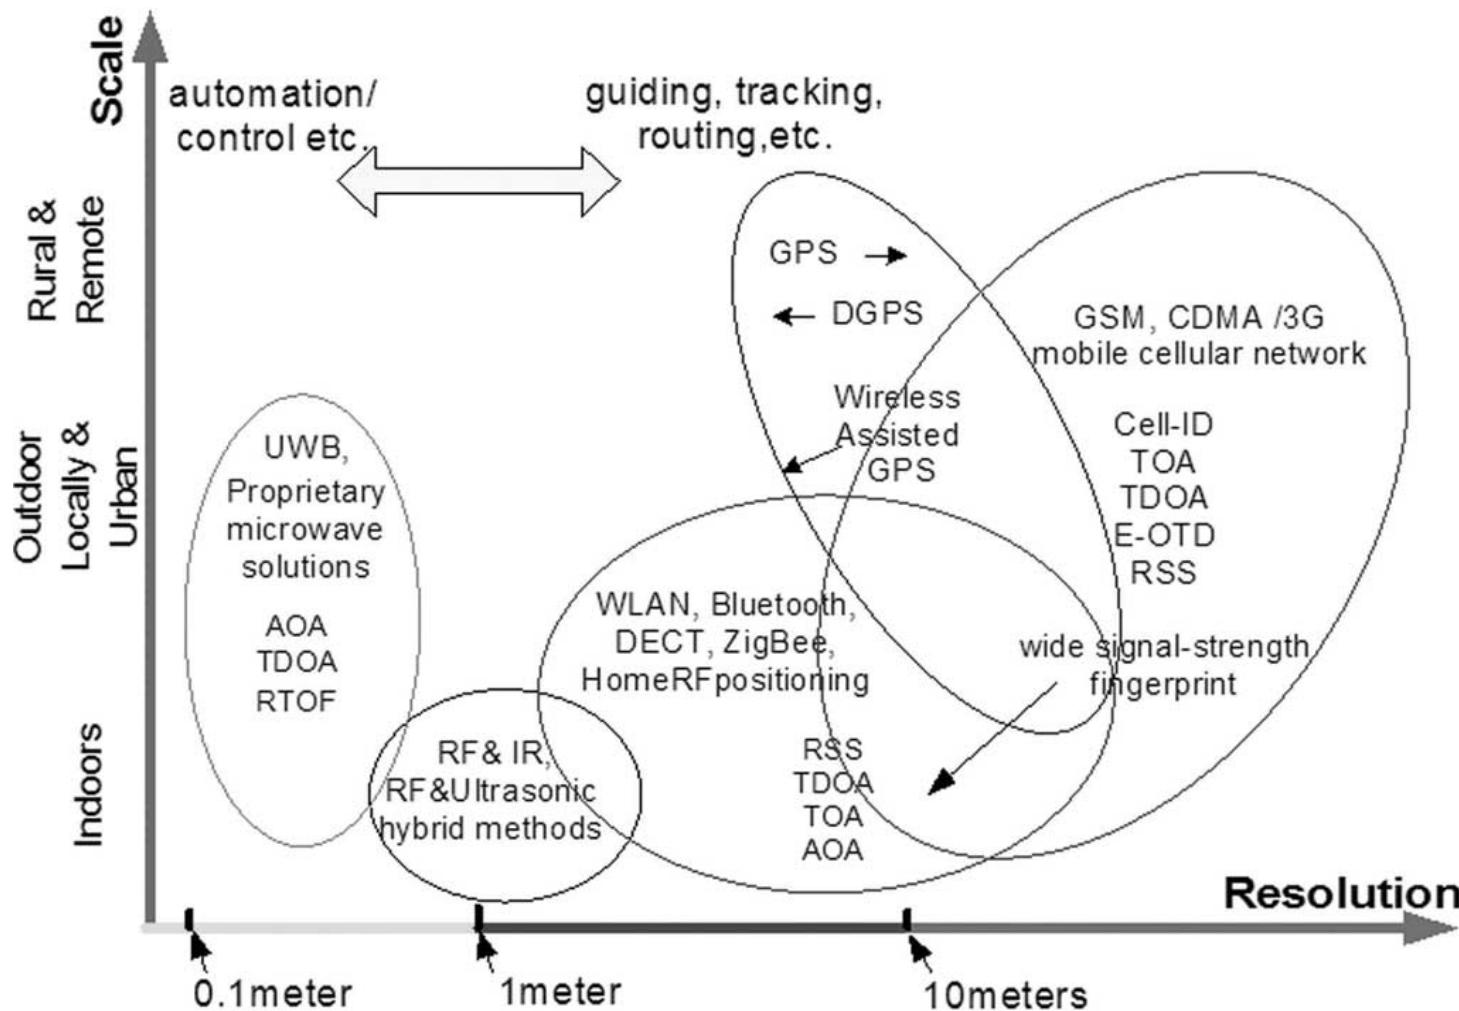


**Continuously find all  
police cars within 3 miles  
of the moving red car**

# Mobile positioning techniques

- **Requirement of location-dependent queries:**
  - ◆ Continuous monitoring of location of the origin of the query, or of the object being queried for – need mobile positioning technologies.
- **Objective of all positioning technologies is to capture the location of a mobile device and convert it into a meaningful X, Y, Z coordinate**
  - ◆ Many technologies are available, deciding which one to use depends on the combination of *accuracy* and *cost*.
  - ◆ When the accuracy increases, so does the cost
  - ◆ Cost is usually shared between the mobile user and wireless carrier

# Mobile positioning techniques



# Location representation

## ■ 2 models

- ◆ Geometric model

specify a location as an n-dimension coordinate

- ◆ Symbolic model

logical, real-world entities describing the location space

e.g., Restaurant A

(250,3000)

{ Mall }

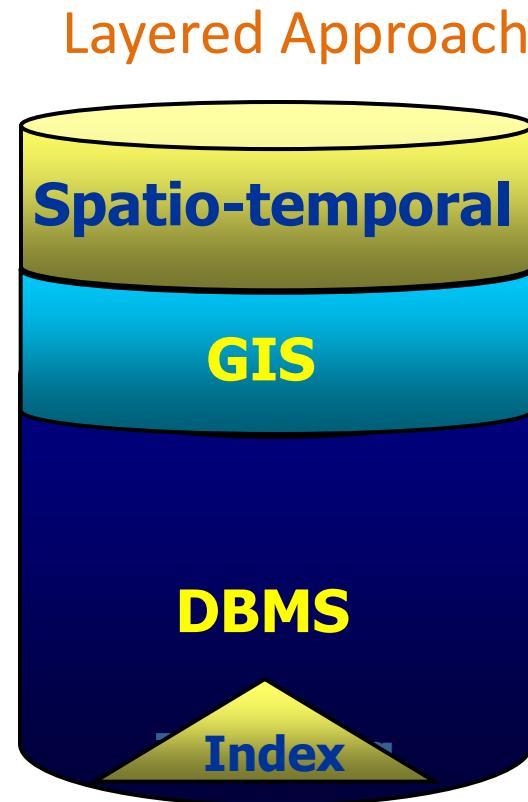
## ■ Spatial data: data representing / associated with locations

- ◆ Space: defined by a set of locations

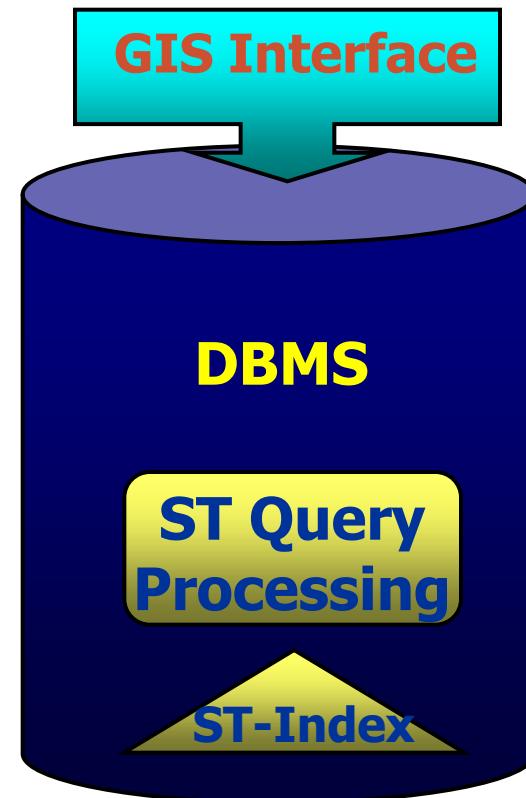
# Location-dependent databases

## Processing of location-dependent queries

- Representation, Indexing and retrieval of location-dependent, spatial, and mobile data



Built-in Approach



# Location-dependent data

- Data whose value is functionally dependent on location.
  - ◆ e.g., city tax, regional weather, hotel rent
- In database, multiple correct values exist - needs *location binding* or *location mapping function*

Report the trucks in my area

```
SELECT O.ID  
FROM Object O  
WHERE O.type = "truck"  
INSIDE Area My-area
```

# Spatial data

## ■ Data about space or entities occupying space.

- ◆ Target at representing the *geometry* of spatial features of the entity, e.g., the location of a road, its orientation.
- ◆ Attribute data represents other information about the spatial features, which are merely regular database data, e.g., the name of a road, its speed limit.

## ■ Spatial objects can be of

- ◆ zero dimension (a *point*),
- ◆ one dimension (a *line*),
- ◆ two dimensions (an *area*)
- ◆ and so on

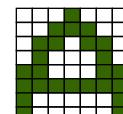
# Spatial data representation

## ■ Spatial features can be *discrete* or *continuous*.

- ◆ Discrete features do not exist between observations but are *individually distinguishable*, e.g., road and land use.
- ◆ Continuous features exist between observations, e.g., elevation, rainfall on land.

## ■ Spatial features can be modeled using *vector data model* or *raster data model*.

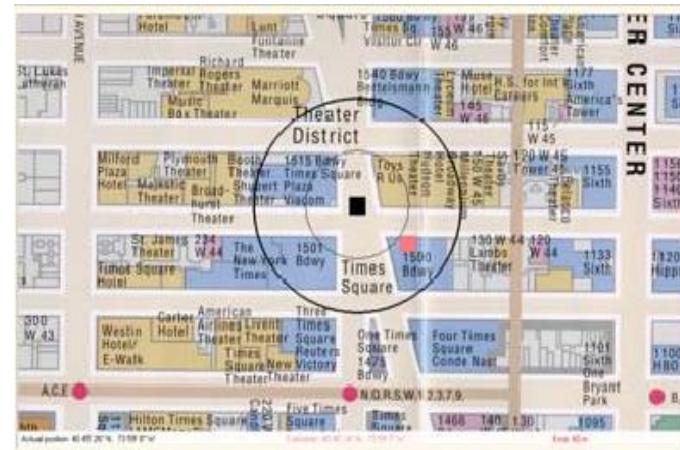
- ◆ **Vector data model:** using points (and their x- and y-coordinates) as basis to model lines, areas and so on.
- ◆ **Raster data model:** using grid to represent spatial feature. Each cell contains a value about the feature at that location.
  - ◆ Generalized raster model is called *tessellation model*, in which the elements can be of *arbitrary shapes*.



# Spatial data representation

■ Geographical Information System (GIS) is the computer-based system to manipulate spatial data

- ◆ A mapping software that translates the low-level position information from the locating infrastructure into meaningful location.
- ◆ building addresses, street layouts, population densities, plethora of other information.



# Spatial queries

- In standard databases, the **comparison operator** is based on **logical and relational operators only**.
- ***Spatial queries*** include **new operators**:
  - ◆ **Containment**: whether a point or an area is contained in another area.
  - ◆ **Overlap**: whether two areas overlap in some part.
  - ◆ **Neighbor**: whether two entities are close to each other.
  - ◆ General **spatial relationship**: e.g., whether a point is to left of a box.
  - ◆ **Spatial join**: joining two relations containing spatial objects, e.g.

```
select L.name  
from Legco L, Company C  
where L.district.area() > 20 and within(C.location, L.district);
```
- General spatial queries include **range query** for objects within a certain area, **nearest neighbor** (object that is nearest to another object), **k-nearest neighbors**.

# Location-dependent vs Spatial queries

- Location-dependent queries are supported by spatial databases and subsume spatial queries.
- Similarity to spatial queries:
  - ◆ Operate on spatial data, i.e. data associated with space occupied by object.
  - ◆ Expressed with spatial query operators, e.g., contains, contained in, intersects, neighboring, north of.
  - ◆ Relying on spatial select and spatial join.
  - ◆ Based on spatial data structures for storage and indexing.
- More than spatial queries:
  - ◆ Client may move in location-dependent queries - location of client may be **part of the query** itself, and result may depend on **direction of movement**.
  - ◆ Data may **not directly contain** location information, but may have multiple values corresponding to different locations.
  - ◆ Requires **dynamic spatial data processing**.
  - ◆ Sometimes **temporal** features are also implied (e.g. gas stations in the next 20 minutes).

# Variety of location-dependent queries



*Continuously report the number of cars in the freeway*

- *Type:* Range query
- *Time:* Present
- *Duration:* Continuous
- *Query:* Stationary
- *Object:* Moving



*What are my nearest McDonalds for the next hour?*

- *Type:* Nearest-Neighbor query
- *Time:* Future
- *Duration:* Continuous
- *Query:* Moving
- *Object:* Stationary



*Send E-coupons to all cars that I am their nearest gas station*

- *Type:* Reverse NN query
- *Time:* Present
- *Duration:* Snapshot
- *Query:* Stationary
- *Object:* Moving



*What was the closest dist. between Taxi A & me yesterday?*

- *Type:* Closest-point query
- *Time:* Past
- *Duration:* Snapshot
- *Query:* Moving
- *Object:* Moving

# Indexing

- Data are stored within a relational database as tuples and tables.
- To answer queries, indexes are built to allow efficient retrieval of data satisfying querying conditions (especially for relational joins).
  - ◆ Allow direct access to tuples satisfying the condition.
  - ◆ Most indexes are built based on hashed table (for point query) and B-tree or B<sup>+</sup>-tree (for range query).
  - ◆ In practice, B<sup>+</sup>-trees are used more often, since they allow sequential access to range data by chaining together leaf nodes.

# Indexing spatial data

- In a spatial database, spatial objects can be stored in their primitive forms
  - e.g., a polygon is represented by the points defining the lines enclosing it or by the polyline.
- Again, indexes must be built to allow effective processing of additional spatial query types.
  - Partition the set of objects into subsets according to their locations
  - Examples include R-tree, R\*-tree, R<sup>+</sup>-tree.
    - ▶ B-tree is based on ordering on a one-dimension key attribute.
    - ▶ Spatial data (e.g. an area) is of two dimensions - R-tree etc. generalize the B-tree and B<sup>+</sup>-tree into their two dimensional equivalence.

# R-Tree

- R-tree is the 2D extension of B<sup>+</sup>-tree.
- In range query, how do you get a set of points or a set of areas contained within a certain bounding box?
  - ◆ A *minimum bounding box* (mbb) is the smallest rectangular region that just completely covers the area.
  - ◆ It is easily computed by taking  $x_1$  and  $x_2$  to be the minimum and maximum value of x-coordinate of all points in the area, and  $y_1$  and  $y_2$  to be the minimum and maximum value of y-coordinate.
- It is not difficult to get a set of points within an area, by generating a range query on the x- and y-coordinate of the corner points of the bounding box (point query).

select p.x, p.y from allPoints p

where x1 <= p.x and p.x <= x2 and y1 <= p.y and p.y <= y2;

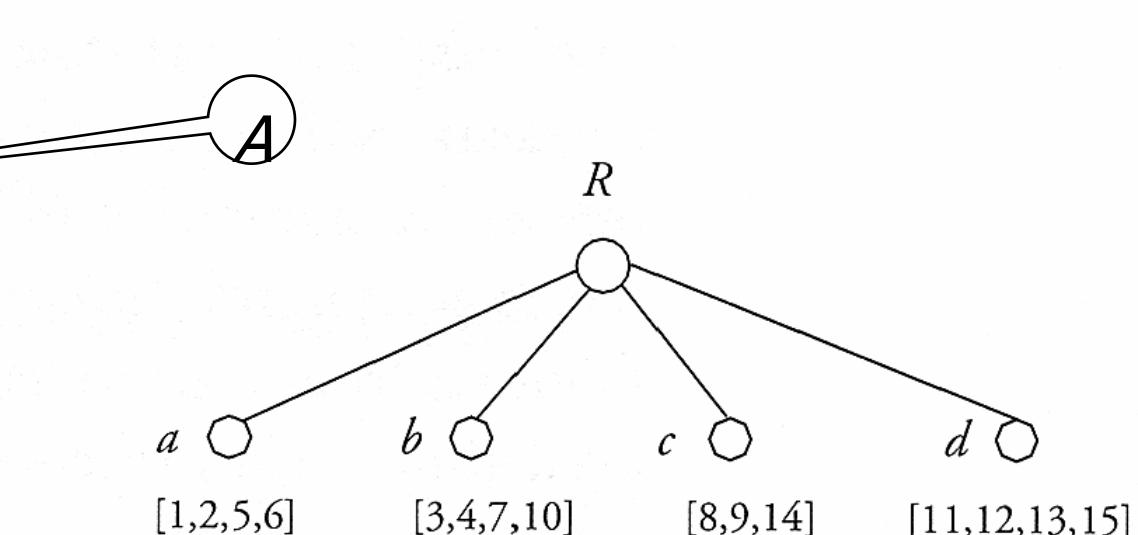
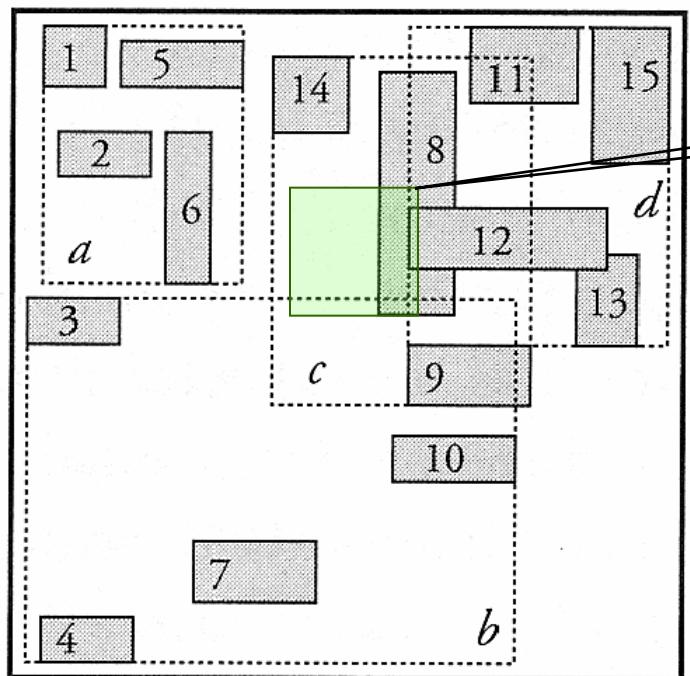
# R-Tree

- In R-Tree, objects (points or rectangles) are clustered into a group, represented by the **mbb** of the group and is considered a node.
- Nodes are considered again as objects for clustering into next level, until the root is reached.
  - ◆ Note that now the **mbb** of interior nodes could overlap.
- For point search, check the **mbb** of each branch.
  - ◆ If the **mbb** of a branch contains the point, search recursively down the sub-branches of the branch.
- For range search, check the **mbb** of each branch.
  - ◆ If the **mbb** of a branch intersects with the query rectangle, search recursively down the sub-branches of the branch.
- Insertion with R-tree is complicated.
- R\*-tree performs some optimizations to R-tree for insertion.

# R-Tree

## ■ An example of an R-tree.

- ◆ How to search for A?



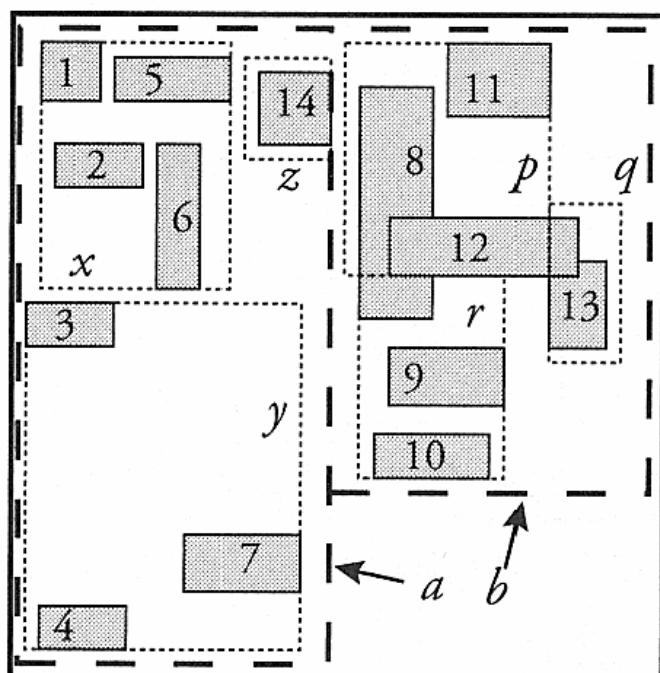
# R<sup>+</sup>-Tree

- Searching in R-tree could be expensive, since multiple branches need to be explored if the mbb of a branch overlaps with the query.
- R<sup>+</sup>-tree is designed to ensure that the mbb for branches at the same level **do not overlap**.
  - ◆ The tradeoff is that the object may have to be duplicated.
- Searching in R<sup>+</sup>-tree is similar to searching in quadtree.
- Compared with R-tree, the storage overhead for R<sup>+</sup>-tree is higher, insertion is more complicated, but searching is much faster.

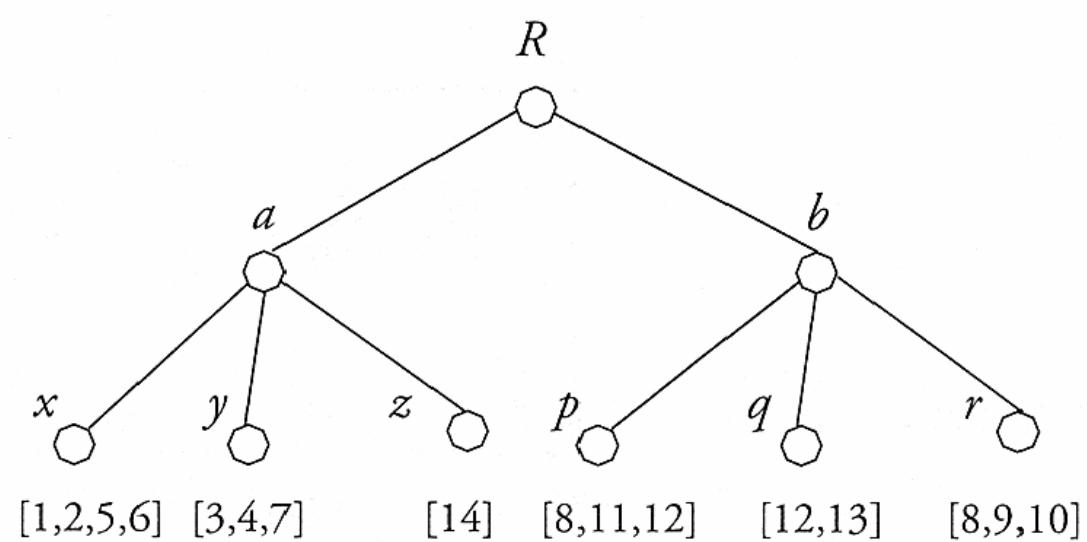
# R<sup>+</sup>-Tree

## ■ An example of an R<sup>+</sup>-tree.

- ◆ Here objects 8 and 12 are duplicated.



$R$



# Indexing moving object

- ***Moving object database*** - spatial database storing locations of moving objects.
  - ◆ E.g. You may want to get the nearest taxi (or the nearest green taxi if you are in Northern N.T.). Database must store the current location of all taxis, since a similar query can be issued at different location.
- Cost of insertion of an object in spatial database is very high. Worse yet, this needs to be done often for a moving object.
- Research is done to search for fast methods to insert and query moving objects, with effective indexing.
- We can trade off precision with update efficiency: update important objects more often but others less often.

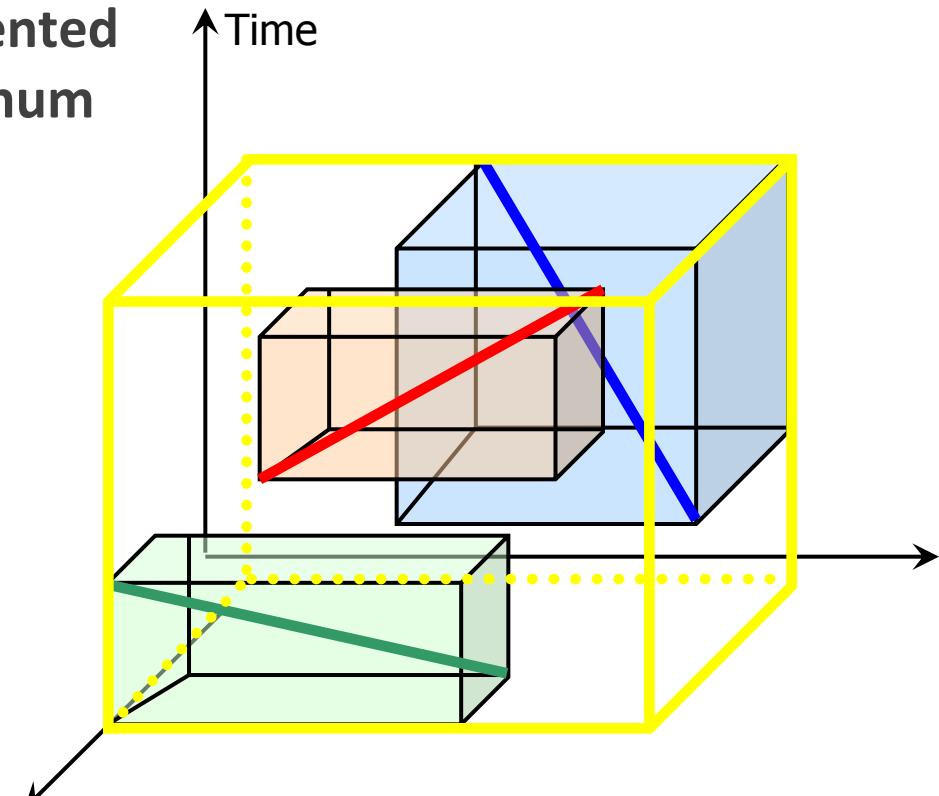
# Indexing time dimension

## ■ Querying Along the Spatial-Temporal Dimension

- ◆ *What was the location of a certain object from 7:00 AM to 10:00 AM yesterday?*
- ◆ *Find all objects that were in a certain area at 7:00 AM yesterday*
- ◆ *Find all objects that were close to each other from 7:00 AM to 8:00 AM yesterday*

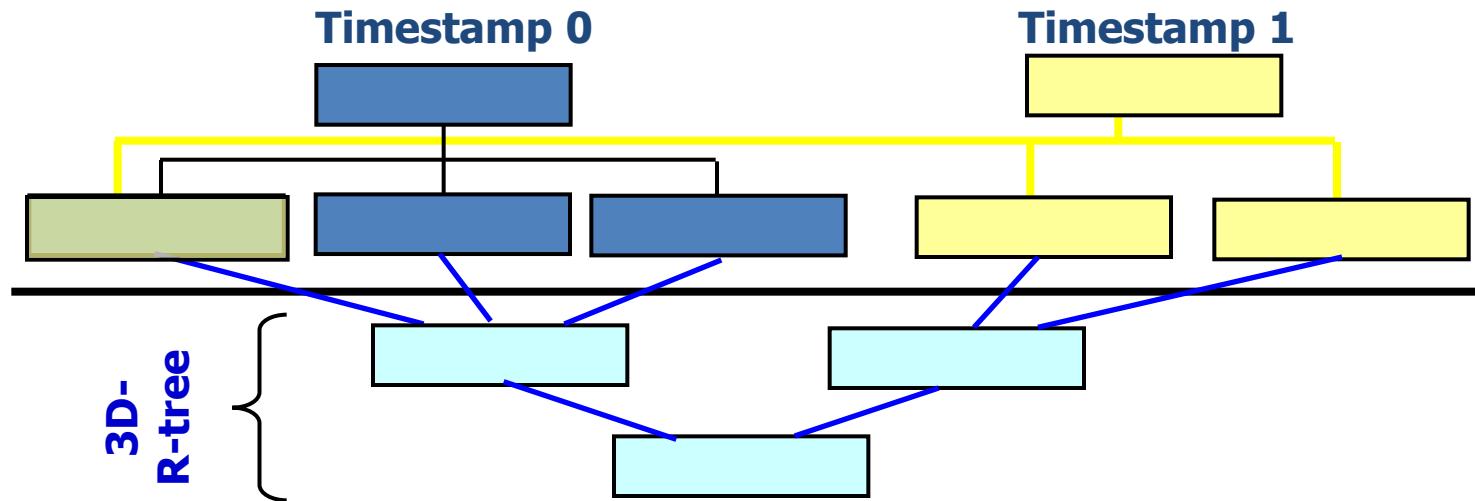
## ■ Historical trajectories are represented by their three-dimensional Minimum Bounding Rectangle (MBR)

- 3D-R-tree is used to index the MBRs
- Technique simple and easy to implement
- Does not scale well
- Does not provide efficient query support



# Indexing time dimension

- Multi-version indexing structures: maintain an R-tree for each time instance
- R-tree nodes that are not changed across consecutive time instances are linked together



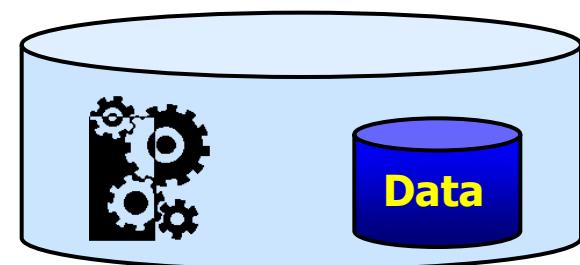
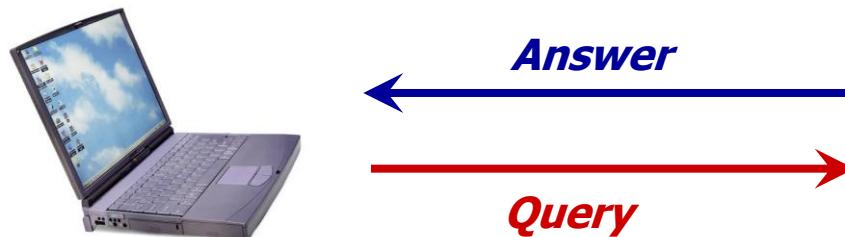
- A multi-version R-tree can be combined with a 3D-R-tree to support interval queries

# Continuous query

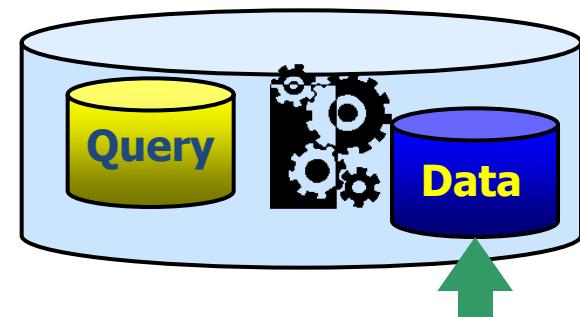
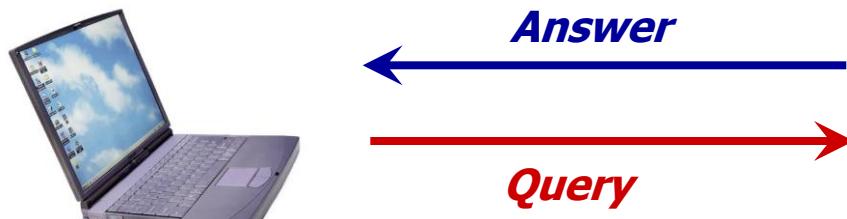
## Query the present – time is always NOW

- ◆ *Find the number of objects in a certain area*
- ◆ *What is the current location of a certain object?*

## Traditional (Snapshot) queries



## Continuous queries



# Continuous query

## ■ *Straightforward Approach*

- ◆ Abstract the continuous queries to a series of snapshot queries evaluated periodically

## ■ *Result Validation*

- ◆ Valid time (t) and Valid region (r)

## ■ *Result Caching*

- ◆ Consecutive evaluations of a continuous query yield similar results
- ◆ Upon evaluation, retrieve more data that can be used later

## ■ *Result Prediction*

- ◆ Based on future trajectory movement, answer can be pre-computed in advance

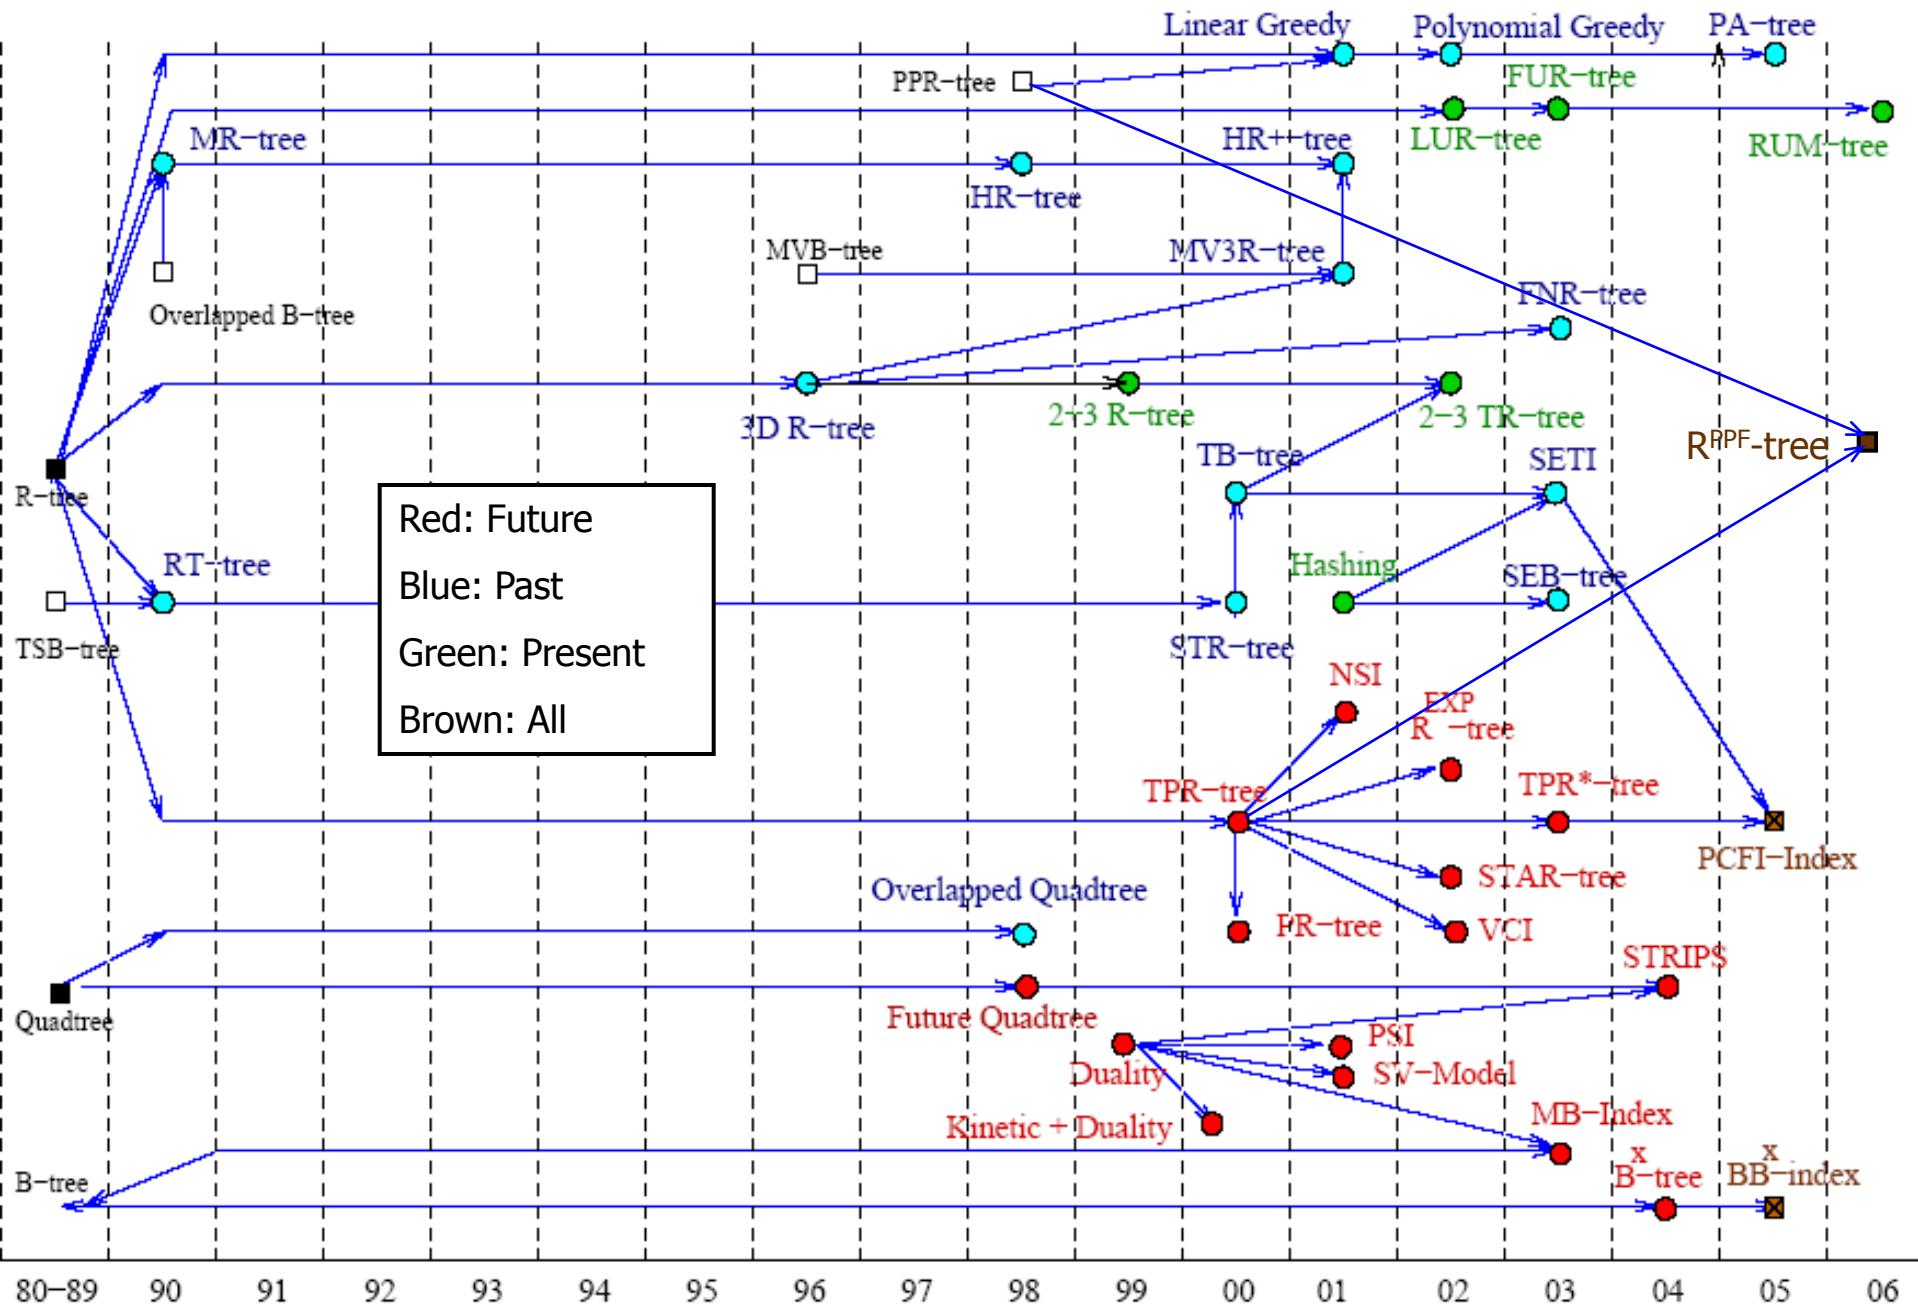
## ■ *Incremental Evaluation*

- ◆ Query is evaluated only once, then only updates of the answer are evaluated – add or delete objects

# Current status of research

- 50+ spatial index structure
- 30+ spatio-temporal indexing structure
- Wide variety of spatio-temporal query processing techniques

# Spatio-temporal access methods



# Open Research Issues

## Location Privacy



*"New technologies can pinpoint your location at any time and place. They promise safety and convenience but threaten privacy and security"*

# Open Research Issues

## Spatio-temporal data mining

- Mining the history → Predicting the future
- Online outlier detection for moving objects
- Suspicious movement in video surveillance
- Analysis of tsunami, hurricanes, or earthquakes
- Phenomena detection and tracking

